

# National University of Computer & Emerging Sciences

## CS 3001 - COMPUTER NETWORKS

### Lecture 14 Chapter 3

10<sup>th</sup> October, 2023

Nauman Moazzam Hayat  
[nauman.moazzam@lhr.nu.edu.pk](mailto:nauman.moazzam@lhr.nu.edu.pk)

Office Hours: 02:30 pm till 06:00 pm (Every Tuesday & Thursday)

# TCP Sender (simplified)

*event: data received from application*

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unACKed segment
  - expiration interval: **TimeoutInterval**

*event: timeout*

- retransmit segment that caused timeout
- restart timer

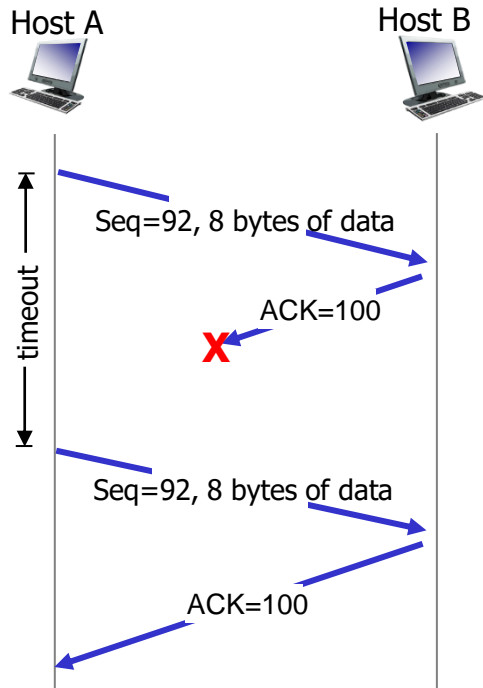
*event: ACK received*

- if ACK acknowledges previously unACKed segments
  - update what is known to be ACKed
  - start timer if there are still unACKed segments

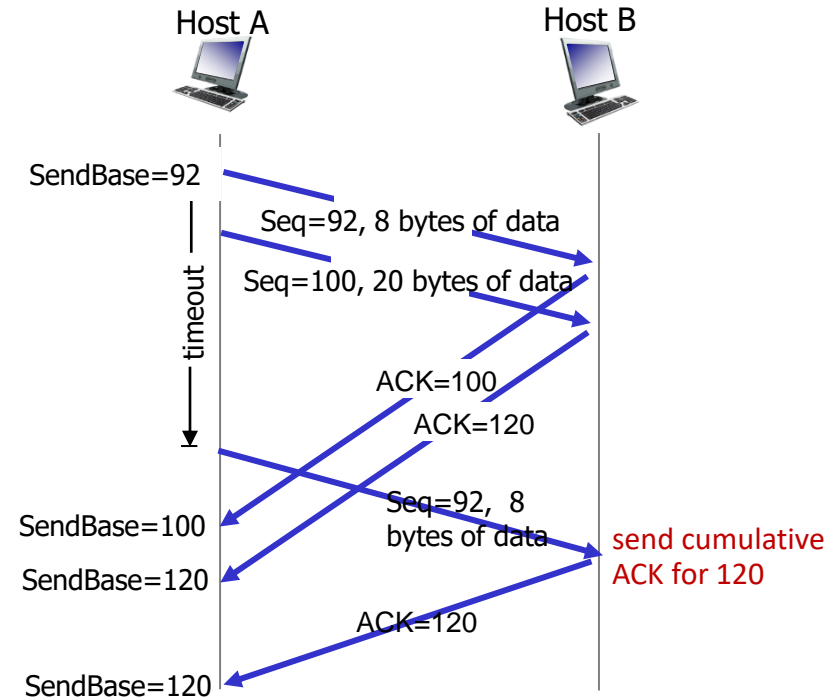
# TCP Receiver: ACK generation [RFC 5681]

<i>Event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

# TCP: retransmission scenarios

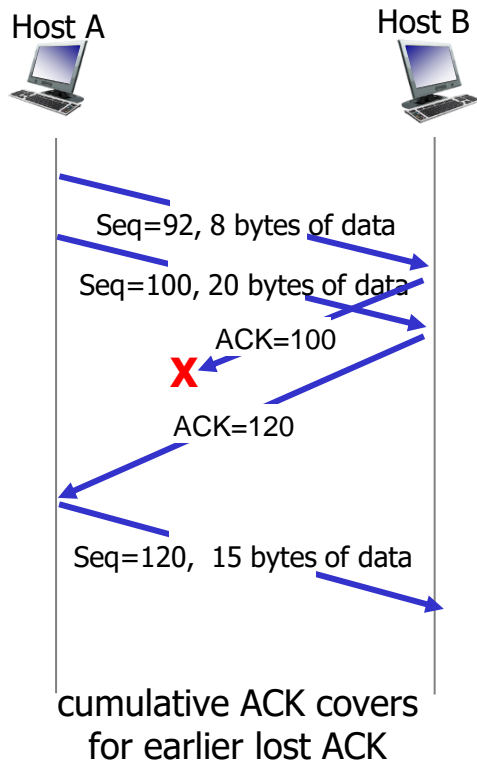


lost ACK scenario



premature timeout

# TCP: retransmission scenarios



# TCP fast retransmit


## *TCP fast retransmit*

if sender receives 3 additional ACKs for same data (“triple duplicate ACKs”), resend unACKed segment with smallest seq #

- likely that unACKed segment lost, so don't wait for timeout



Receipt of three duplicate ACKs indicates 3 segments received after a missing segment – lost segment is likely. So retransmit!




# Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- **Connection-oriented transport: TCP**
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- Principles of congestion control
- TCP congestion control



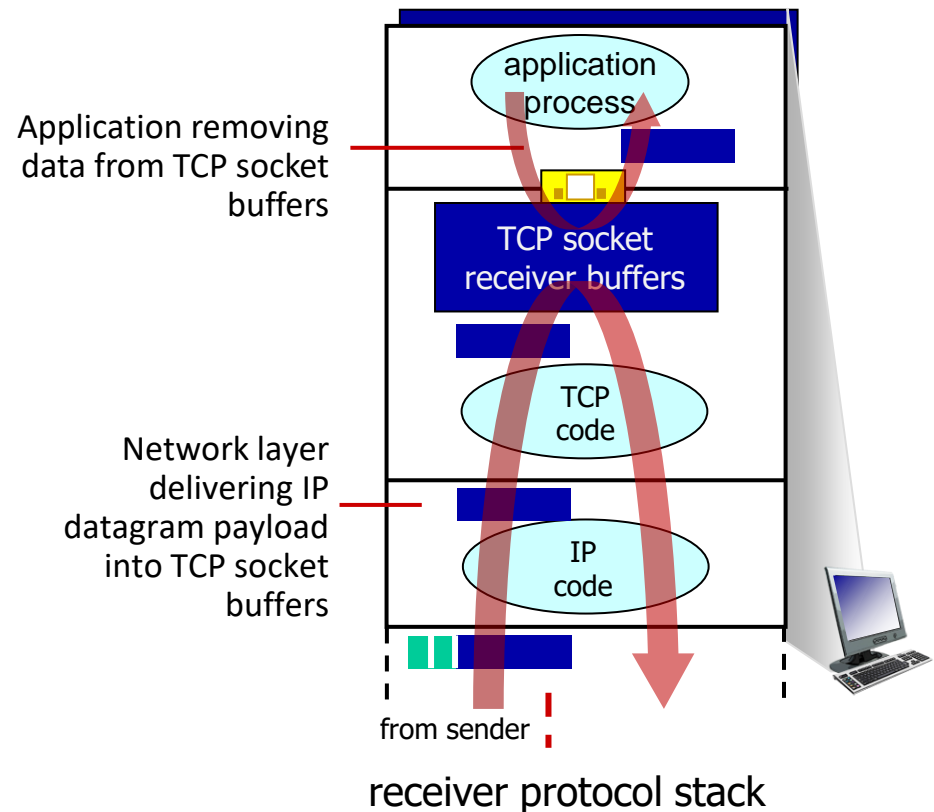
# Flow Control vs Congestion Control (Courtesy Rutgers University)





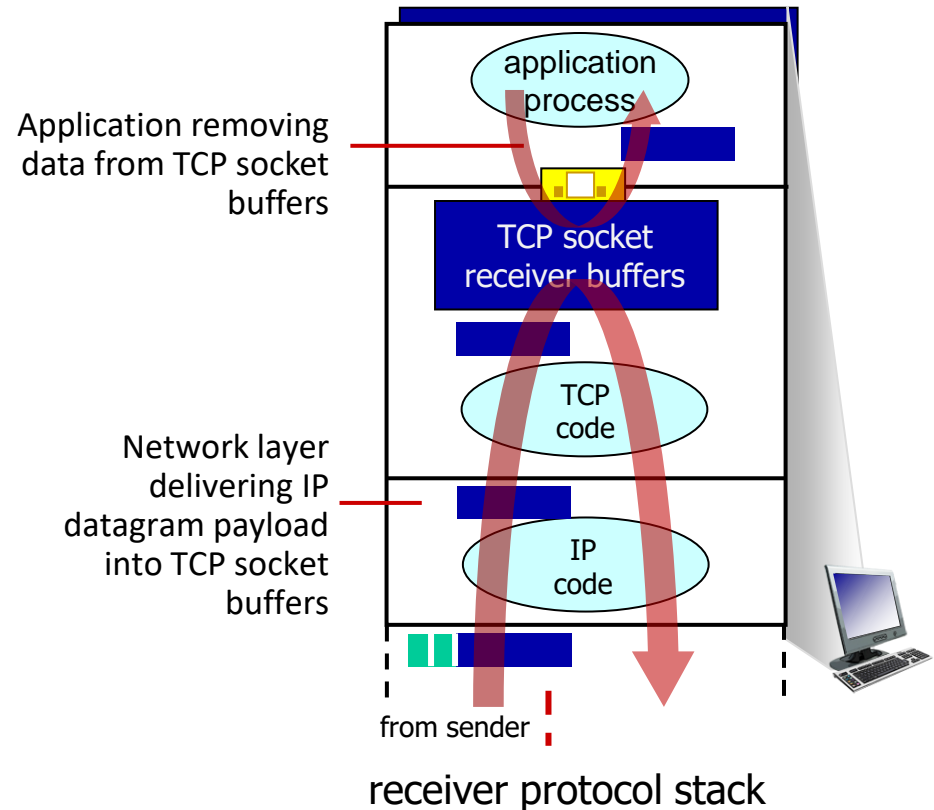
# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



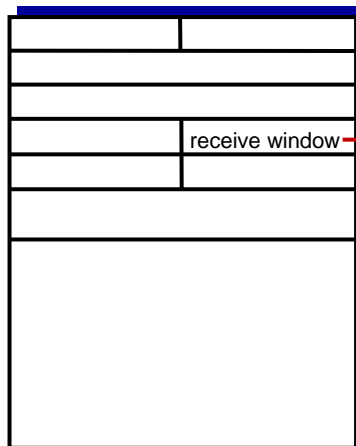
# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



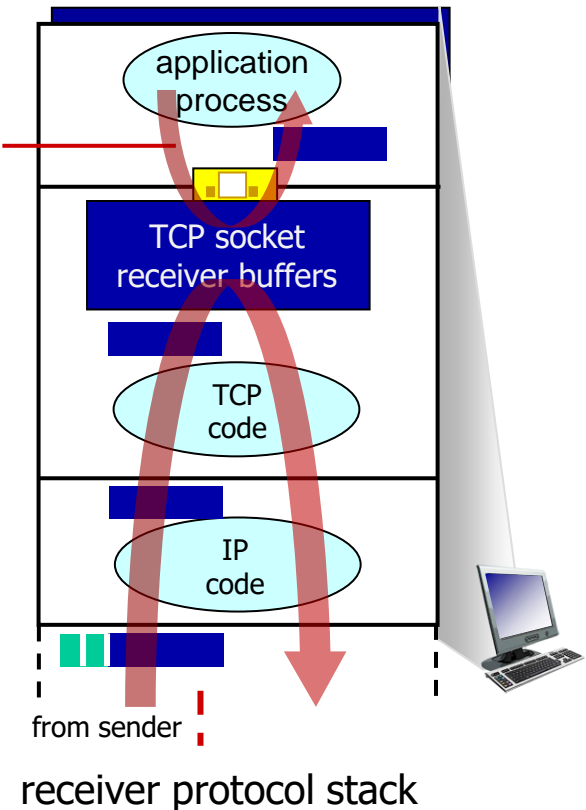
# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?



flow control: # bytes receiver willing to accept

Application removing data from TCP socket buffers

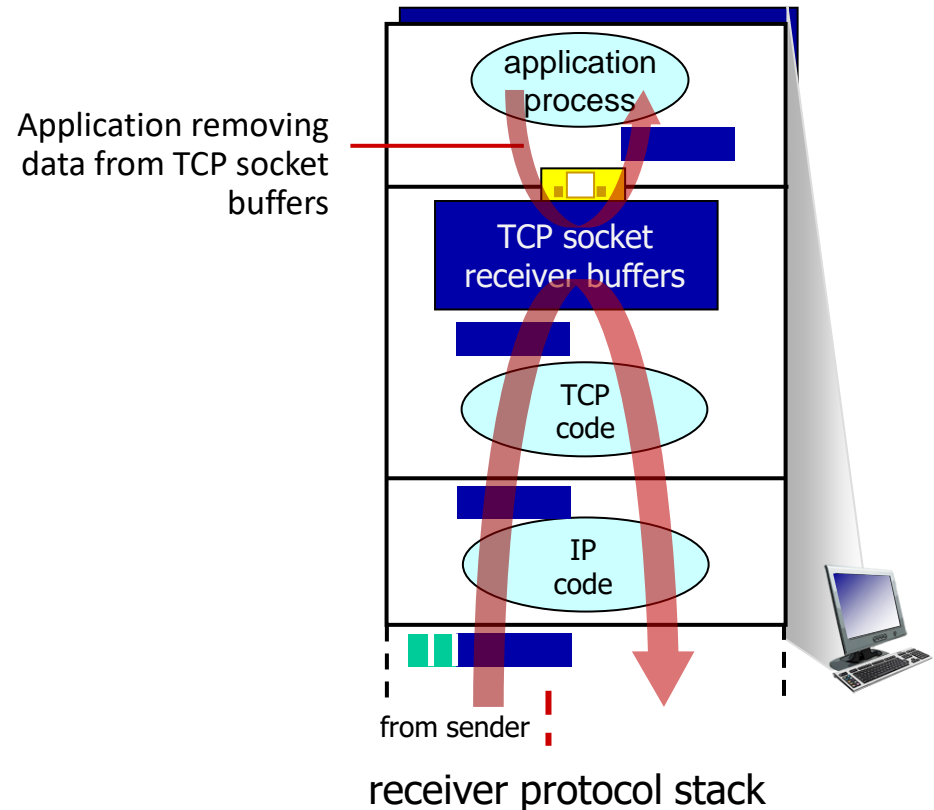


# TCP flow control

Q: What happens if network layer delivers data faster than application layer removes data from socket buffers?

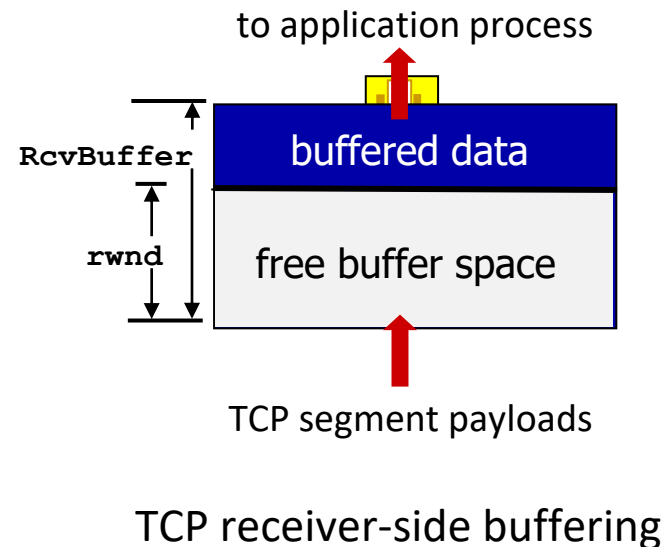
## flow control

receiver controls sender, so sender won't overflow receiver's buffer by transmitting too much, too fast



# TCP flow control


- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems auto-adjust **RcvBuffer**
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow



# TCP flow control

- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems auto-adjust **RcvBuffer** (**i.e. dynamic**)
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow
- (**Assumption: TCP receiver discards out-of-order segments**)

flow control: # bytes receiver willing to accept



TCP segment format

# TCP Flow Control - Implementation

- ❖ TCP provides flow control by having the sender maintain a dynamic variable called the receive window (**rwnd**)

(Since TCP is not permitted to overflow the allocated buffer, **we must have**)

**$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{RcvBuffer}$**  (where:

- **RcvBuffer** is the size of the buffer allocated at the receiver for this connection,
- **LastByteRcvd** is the number of the last byte in the data stream received from the network and placed in the receive buffer, &
- **LastByteRead** is the number of the last byte in the data stream read by the application process from the receive buffer), **thus**

**$\text{rwnd} = \text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$**

- ❖ While the receiver is keeping track of many variables as seen above, the sender is keeping track of primarily two variables, (i.e. **LastByteSent & LastByteAcked**)

**$\text{LastByteSent} - \text{LastByteAcked} \leq \text{rwnd}$**  (i.e. the **UnAckedData**.)

- By maintaining this throughout the connection's life, i.e. keeping the **UnAckedData** less than or equal to the **rwnd**, the sender ensures it doesn't overflow the buffer at the receiver

# TCP Flow Control - Issue

**Issue:** One minor technical problem with this scheme.

- To see this, suppose Host B's (receiver) receive buffer becomes full so that  $\text{rwnd} = 0$ .
- After advertising  $\text{rwnd} = 0$  to Host A (sender), also suppose that B has nothing to send to A.
- As the application process at B empties the buffer, TCP does not send new segments with new  $\text{rwnd}$  values to Host A (indeed, TCP sends a segment to Host A only if it has data to send or if it has an acknowledgment to send)

Therefore, Host A is never informed that some space has opened up in Host B's receive buffer (i.e. **Host A is blocked and can transmit no more data!**)


- To **solve** this problem, the TCP specification requires Host A to continue to send segments with one data byte when B's receive window is zero.
- These segments will be acknowledged by the receiver.
- Eventually the buffer will begin to empty and the acknowledgments will contain a non-zero  $\text{rwnd}$  value.




# TCP connection management

before exchanging data, sender/receiver “handshake”:

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters (e.g., starting seq #s)



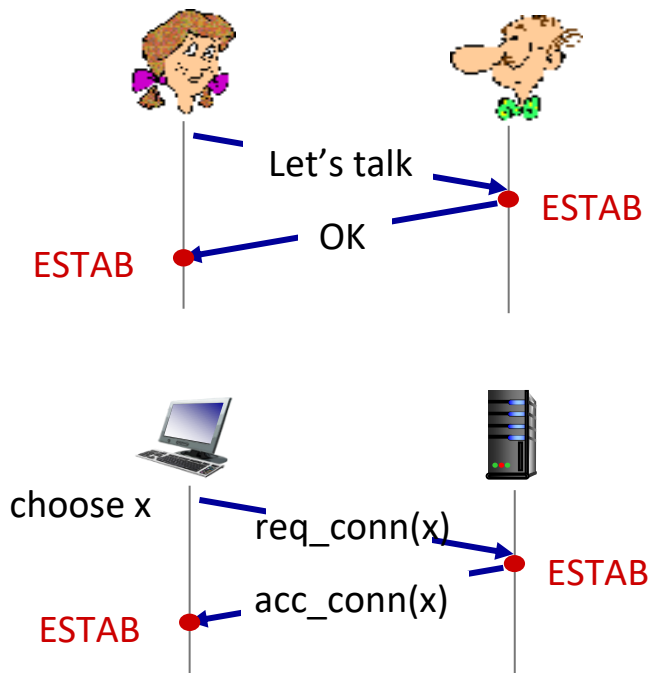
```
Socket clientSocket =  
    newSocket("hostname", "port number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

# Agreeing to establish a connection


2-way handshake:



Q: will 2-way handshake always work in network?

- variable delays
- retransmitted messages (e.g. req\_conn(x)) due to message loss
- message reordering
- can't "see" other side

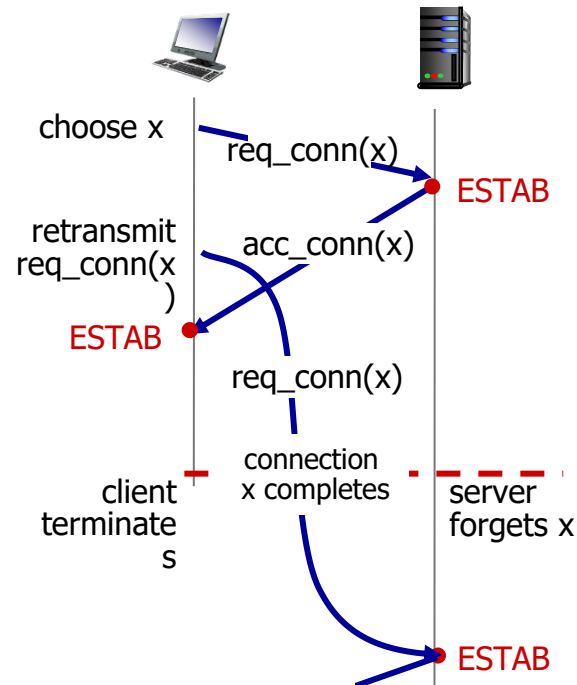
# 2-way handshake scenarios




No problem!

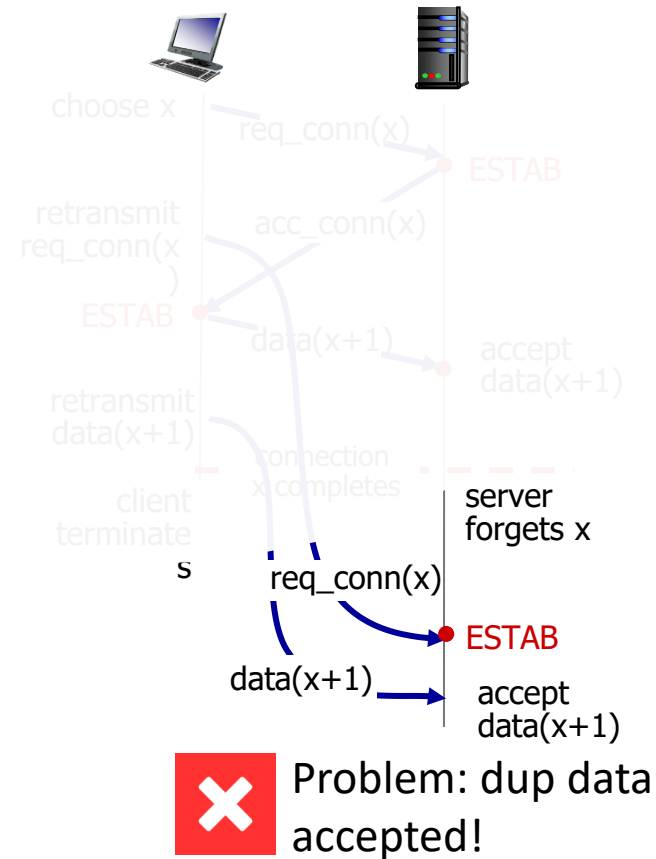


## 2-way handshake scenarios



 Problem: half open connection! (no client)

# 2-way handshake scenarios



# TCP 3-way handshake

## Client state

```
clientSocket = socket(AF_INET, SOCK_STREAM)
```

LISTEN

```
clientSocket.connect((serverName, serverPort))
```

SYNSENT

ESTAB

choose init seq num, x  
send TCP SYN msg

received SYNACK(x)  
indicates server is live;  
send ACK for SYNACK;  
this segment may contain  
client-to-server data



SYNbit=1, Seq=x

SYNbit=1, Seq=y  
ACKbit=1; ACKnum=x+1

ACKbit=1, ACKnum=y+1



choose init seq num, y  
send TCP SYNACK  
msg, acking SYN

received ACK(y)  
indicates client is live

## Server state

```
serverSocket = socket(AF_INET, SOCK_STREAM)  
serverSocket.bind('', serverPort)  
serverSocket.listen(1)  
connectionSocket, addr = serverSocket.accept()
```

LISTEN

SYN RCVD

ESTAB

# A human 3-way handshake protocol



# Closing a TCP connection

- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled



# Assignment # 3 (Chapter - 3)

- *3<sup>rd</sup> Assignment will be uploaded on Google Classroom on Thursday, 12<sup>th</sup> October, 2023, in the Stream - Announcement Section*
- *Due Date: Tuesday, 17<sup>th</sup> October, 2023 (Handwritten solutions to be submitted during the lecture)*
- *Please read **all the instructions** carefully in the uploaded Assignment document, follow & submit accordingly*

## Quiz # 3 (Chapter - 3)

- *On: Thursday, 19<sup>th</sup> October, 2023 (During the lecture)*
- *Quiz to be taken during own section class only*